

Lab4.R

rstudio-user

2021-02-08

#1. Read dataset


```
data <- read.csv('marketing.csv')  
print(data)
```

##	Month	Spend	Sales
## 1	1	1000	9914
## 2	2	4000	40487
## 3	3	5000	54324
## 4	4	4500	50044
## 5	5	3000	34719
## 6	6	4000	42551
## 7	7	9000	94871
## 8	8	11000	118914
## 9	9	15000	158484
## 10	10	12000	131348
## 11	11	7000	78504
## 12	12	3000	36284

#2. Use Scatter Plot to visualize the Relationship

```
scatter.smooth(x=data$Spend, y=data$Sales, main="Sales~Spend", xlab="Spend", ylab="Sales")
```

Sales~Spend




#With the scatter plot, we can easily see that the data has a positive strong correlation.

#3. Using BoxPlot to check for Outliers


```
boxplot(data$Spend, main="Outliers in Spend")
```

Outliers in Spend



```
boxplot(data$Sales, main="Outliers in Sales")
```

Outliers in Sales



#There are no outliers in either of the columns as there are no data point outside the plot range.

#4. Using Density Plot To Check If Response Variable Is Close To Normal

```
library(e1071)
```



```
par(mfrow=c(1, 2))
```

```
plot(density(data$Spend), main="Density : Spend", xlab="Spend", ylab="Frequency", sub=paste("Skewness:"
```

```
polyon(density(data$Spend), col="blue")
```

```
plot(density(data$Sales), main="Density : Spend", xlab="Sales", ylab="Frequency", sub=paste("Skewness:"
```

```
polyon(density(data$Sales), col="green")
```



#5. Check the Correlation Analysis

```
cor(data$Spend, data$Sales)
```

```
## [1] 0.9988322
```

#Since the correlation of the two variables is above 0.99, approx 1, we can say that they have a strong

#6. Build the Linear Regression Model

```
model <- lm(Sales~Spend, data=data)
```

#7. Using p-value Check For Statistical Significance

```
modelsummary <- summary(model)
modelcoeffs <- modelsummary$coefficients
est <- modelcoeffs['Spend', 'Estimate']
stderr <- modelcoeffs['Spend', 'Std. Error']
tval <- est/stderr
pval <- 2*pt(-abs(tval), df=length(data$Sales)-1)
print(paste("p-value:", pval))
```

```
## [1] "p-value: 1.32941244094728e-15"
```

#8. Capture the summary of the linear model

```
modelsummary
```

```
##
```

```
## Call:
```

```
## lm(formula = Sales ~ Spend, data = data)
```



```
##
```



```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -3385 -2097 258 1726 3034
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1383.4714 1255.2404 1.102 0.296
## Spend 10.6222 0.1625 65.378 1.71e-14 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2313 on 10 degrees of freedom
## Multiple R-squared: 0.9977, Adjusted R-squared: 0.9974
## F-statistic: 4274 on 1 and 10 DF, p-value: 1.707e-14
```

#9. Also perform the Linear Diagnostics for the given data set(Hint: plot(lmmodel))
`plot(model)`





```
#10. Create the training and test data (70:30)
```

```
set.seed(11)
```

```
indexes = sample(nrow(data), nrow(data)*.7)
```

```
train <- data[indexes,]
```

```
test <- data[-indexes,]
```

```
#11. Fit the model on training data and predict sales on test data
```

```
model <- lm(Sales~Spend, data=train)
```

```
predict(model)
```

```
##          10          2          8          9          12          1          5          4
## 128959.82  43974.49 118336.65 160829.32 33351.33 12104.99 33351.33 49286.07
```

```
#We can see the values predicted by the model. On comparing a value at an index with the observed value
```

```
#12. Review the diagnostic measures
```

```
plot(model)
```

